

[70240413 Statistical Machine Learning, Spring, 2017]

# **Deep Learning** **(deep neural nets)**

**Jun Zhu**

`dcszj@mail.tsinghua.edu.cn`

`http://bigml.cs.tsinghua.edu.cn/~jun`

State Key Lab of Intelligent Technology & Systems

Tsinghua University


March 21, 2017

# Why going deep?

- ◆ Data are often high-dimensional.
- ◆ There is a huge amount of **structure** in the data, but the structure is too complicated to be represented by a simple model.
- ◆ Insufficient depth can require more **computational elements** than architectures whose depth matches the task.
- ◆ Deep nets provide simpler but more descriptive models of many problems.

# Resolution in Image Classification

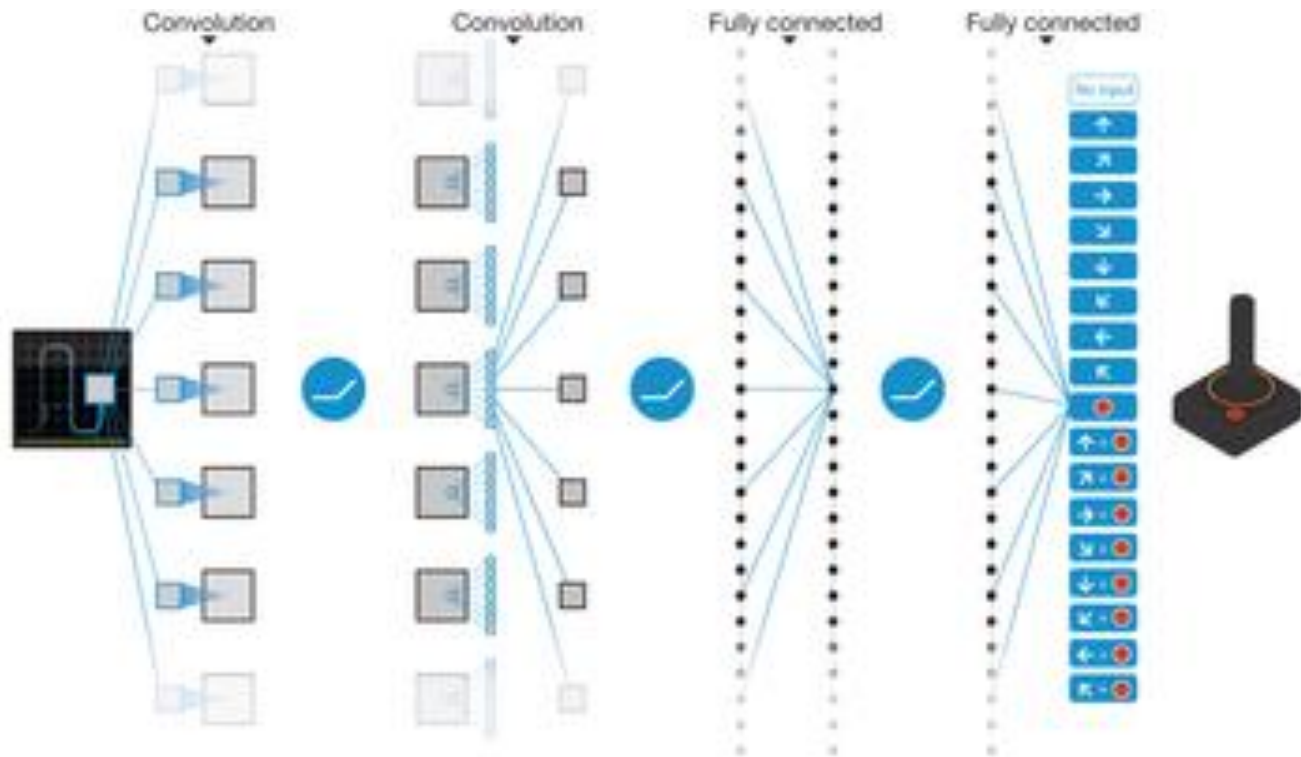

◆ ImageNet Large-Scale Visual Recognition Challenge (ILSVRC)






# Human-Level Control via Deep RL

- ◆ Deep Q-network with human-level performance on Atari games

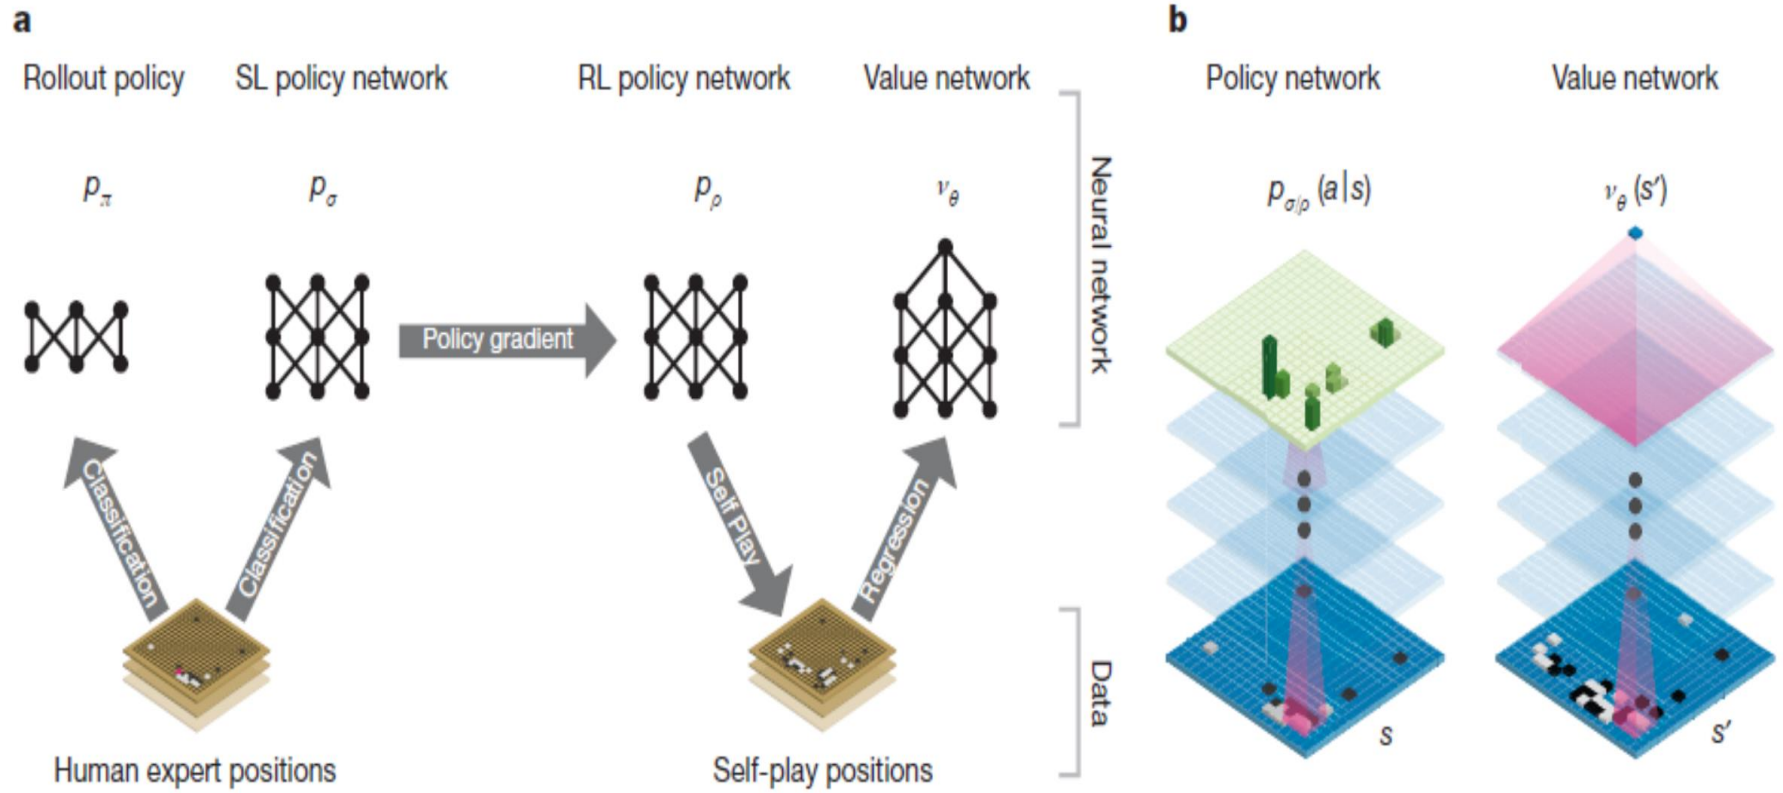


[Mnih et al., Nature 518, 529–533, 2015]



# AlphaGo


## ◆ Neural network training pipeline and architecture






# AlphaGo

## ◆ Monte Carlo tree search






# AlphaGo



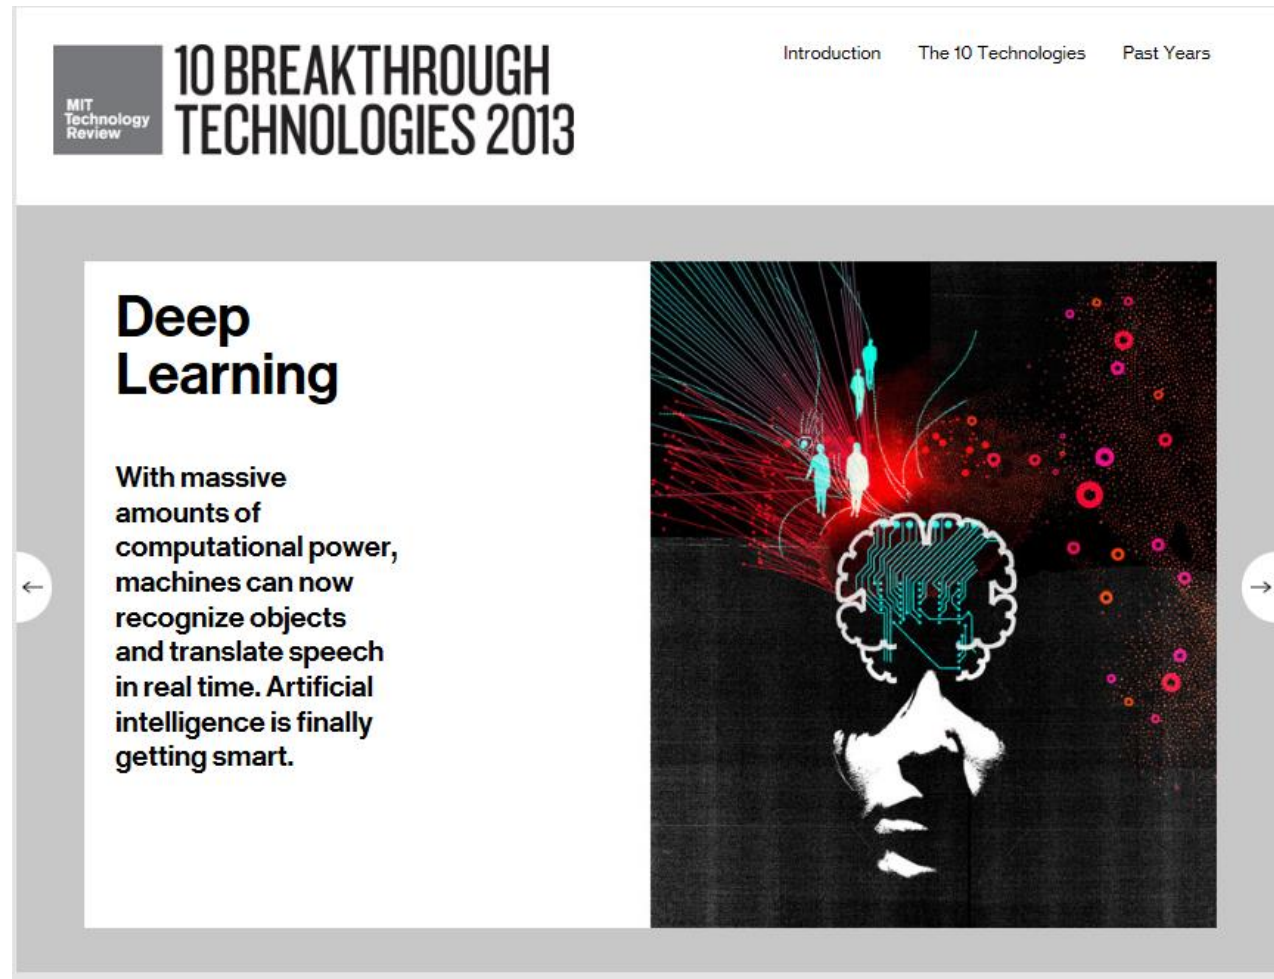


◆ Tencent FineArt





# MIT 10 Breakthrough Tech 2013



<http://www.technologyreview.com/featuredstory/513696/deep-learning/>

# Deep Learning in industry



Driverless car



Face identification



Speech recognition



Web search

...



...



# History of neural networks



Pitts



McCulloch



Rosenblatt



Minsky



Papert



Ackley




Hinton




Sejnowski

1943



1960




1969



Perceptrons

1985



Boltzmann Machine



# History of neural networks




Smolensky



Hinton

Hinton et al.

1986




Harmoniums  
(Restricted Boltzmann Machine)

2002



Contrastive  
Divergence

2006






Deep Belief  
Networks

# Deep Learning Models



# How the human brain learns?




The business end of this is made of lots of these joined in networks like this

Much of our own “computations” are performed in/by this network

Learning occurs by changing the effectiveness of the synapses so that the influence of one neuron on another changes

# How the human brain learns?

## ◆ A typical neuron




## ◆ Learning occurs by changing the effectiveness of the synapses so that the influence of one neuron on another changes





# Model of a neuron





$$v_k = \sum_{j=1}^m w_{kj} x_j + b_k$$

$$y_k = \psi(v_k)$$

# Activation function


◆ Threshold function & piecewise linear function:



◆ Sigmoid function

$$\psi_{\alpha}(v) = \frac{1}{1 + \exp(-\alpha v)}$$

$a \rightarrow \infty$  : step function




# Activation function with negative values

◆ Threshold function & piecewise linear function:

$$\text{sgn}(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ -1 & \text{if } x < 0 \end{cases}$$


◆ Hyperbolic tangent function

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



# McCulloch & Pitts's Artificial Neuron


- ◆ The first model of artificial neurons in 1943
  - Activation function: a threshold function





$$y_j = \text{sgn} \left( \sum_i w_{ij} x_i - \theta_j \right)$$

# Network Architecture

## ◆ Feedforward networks




## ◆ Recurrent networks





# Learning Paradigms

- ◆ Supervised Learning (learning with a teacher)
  - For example, classification: learns a separation plane



# Learning Paradigms

- ◆ Unsupervised learning (learning without a teacher)
  - Example: clustering



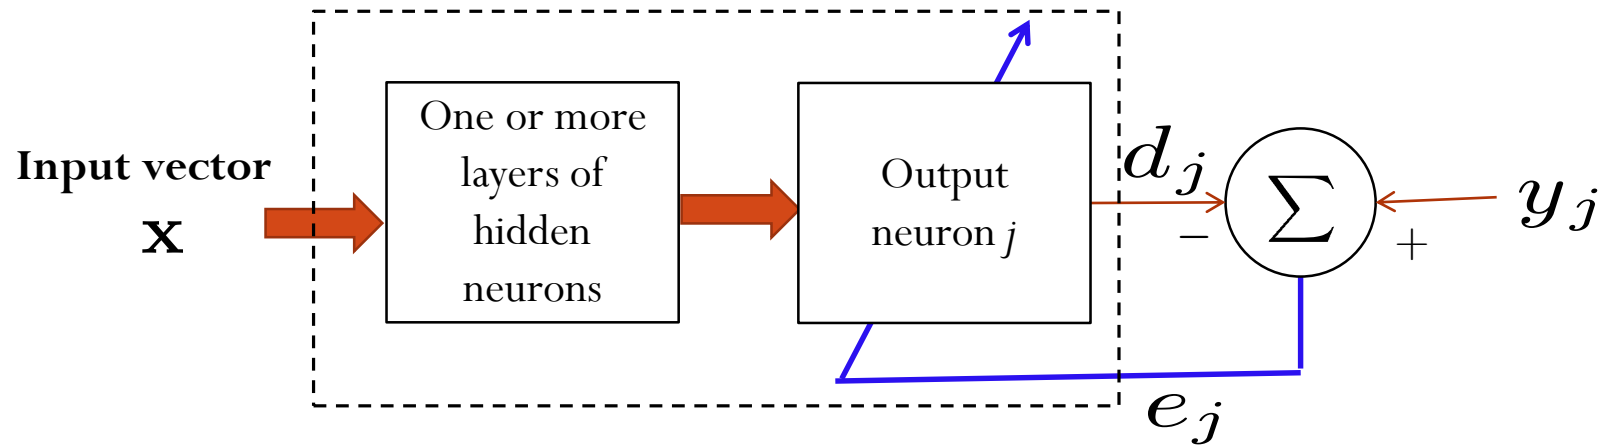
# Learning Rules

- ◆ Error-correction learning
- ◆ Competitive learning
- ◆ Hebbian learning
- ◆ Boltzmann learning
- ◆ Memory-based learning
  - Nearest neighbor, radial-basis function network



# Error-correction learning

◆ The generic paradigm:



□ Error signal:


$$e_j = y_j - d_j$$

□ Learning objective:

$$\min_{\mathbf{w}} R(\mathbf{w}; \mathbf{x}) := \frac{1}{2} \sum_j e_j^2$$

## Example: Perceptron

- ◆ One-layer feedforward network based on error-correction learning (no hidden layer):



- Current output (at iteration  $t$ ):

$$d_j = (\mathbf{w}_t^j)^\top \mathbf{x}$$

- Update rule (*exercise?*):

$$\mathbf{w}_{t+1}^j = \mathbf{w}_t^j + \eta(y_j - d_j)\mathbf{x}$$

# Perceptron for classification

- ◆ Consider a single output neuron

- ◆ Binary labels:

$$y \in \{+1, -1\}$$

- ◆ Output function:

$$d = \text{sgn} \left( \mathbf{w}_t^\top \mathbf{x} \right)$$

- ◆ Apply the error-correction learning rule, we get ... (next slide)



# Perceptron for Classification

◆ Set  $\mathbf{w}_1 = 0$  and  $t=1$ ; scale all examples to have length 1  
(doesn't affect which side of the plane they are on)

◆ Given example  $\mathbf{x}$ , predict positive *iff*


$$\mathbf{w}_t^\top \mathbf{x} > 0$$

◆ If a mistake, update as follows

□ Mistake on positive:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \eta_t \mathbf{x}$

□ Mistake on negative:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \eta_t \mathbf{x}$

$t \leftarrow t + 1$



# Convergence Theorem

- ◆ For linearly separable case, the perceptron algorithm will converge in a finite number of steps



# Mistake Bound

## ◆ Theorem:

- Let  $\mathcal{S}$  be a sequence of labeled examples consistent with a linear threshold function  $\mathbf{w}_*^\top \mathbf{x} > 0$ , where  $\mathbf{w}_*$  is a unit-length vector.
- The number of mistakes made by the online Perceptron algorithm is at most  $(1/\gamma)^2$ , where


$$\gamma = \min_{\mathbf{x} \in \mathcal{S}} \frac{|\mathbf{w}_*^\top \mathbf{x}|}{\|\mathbf{x}\|}$$

- i.e.: if we scale examples to have length 1, then  $\gamma$  is the minimum distance of any example to the plane  $\mathbf{w}_*^\top \mathbf{x} = 0$
- $\gamma$  is often called the “margin” of  $\mathbf{w}_*$ ; the quantity  $\frac{\mathbf{w}_*^\top \mathbf{x}}{\|\mathbf{x}\|}$  is the cosine of the angle between  $\mathbf{x}$  and  $\mathbf{w}_*$

# Deep Nets

- ◆ Deep neural networks
  - Multi-layer Perceptron
  - CNN
  - Deep recurrent nets
- ◆ Deep generative models
  - Auto-encoder
  - RBM
  - Deep belief nets

# XOR Problem

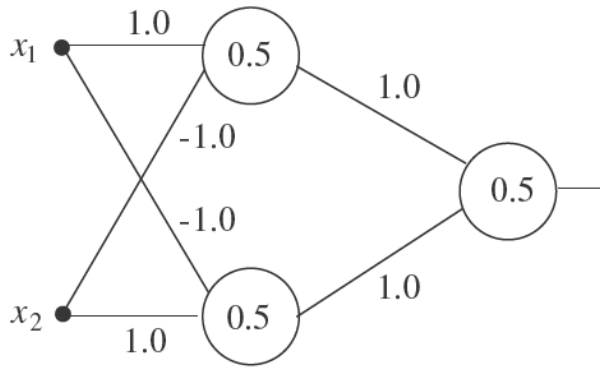


◆ Single-layer perceptron can't solve the problem

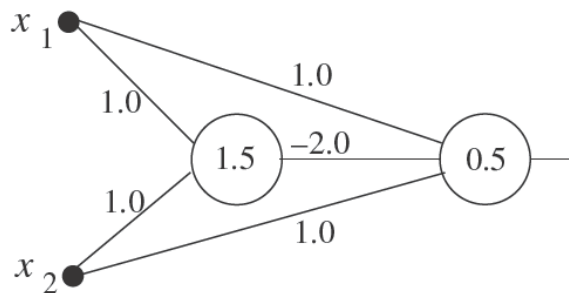


# XOR Problem

- ◆ A network with 1-layer of 2 neurons works for XOR:
  - threshold activation function



- Many alternative networks exist (not layered)



# Multilayer Perceptrons

- ◆ Computational limitations of single-layer Perceptron by Minsky & Papert (1969)
  
- ◆ Multilayer Perceptrons:
  - Multilayer feedforward networks with an error-correction learning algorithm, known as error *back-propagation*
  
  - A generalization of single-layer perceptron to allow nonlinearity

# Backpropagation

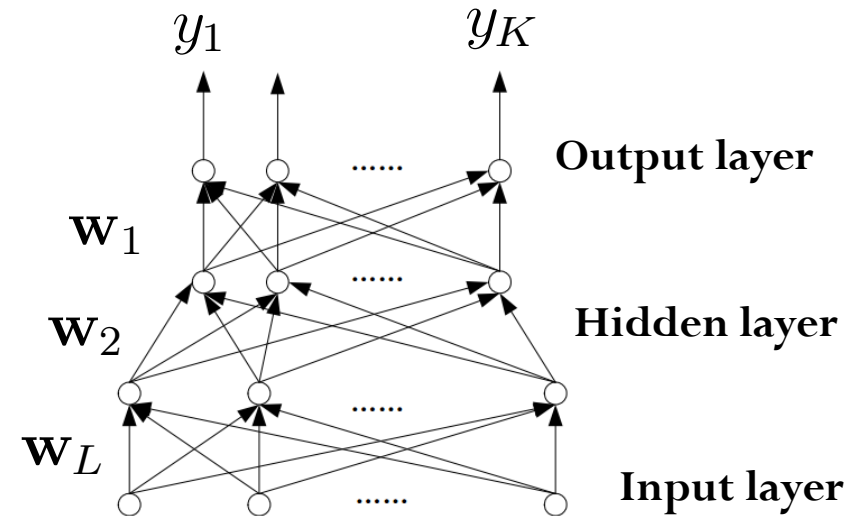
- ◆ Learning as loss minimization

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{2} \sum_j e_j^2(\mathbf{x})$$

$$e_j = y_j - d_j$$

- ◆ Learning with gradient descent

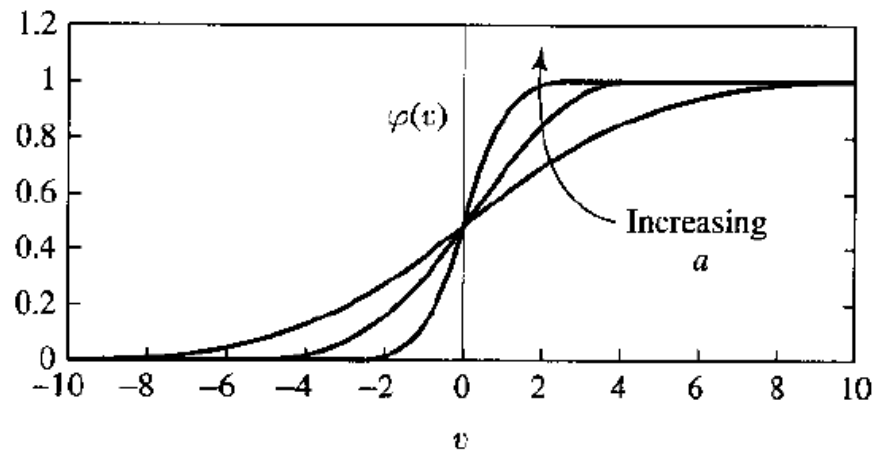
$$\mathbf{w}_{t+1} = \mathbf{w}_t - \lambda_t \nabla R(\mathbf{w}; \mathcal{D})$$



# Backpropagation

- ◆ Step function in perceptrons is non-differentiable
- ◆ Differentiable activation functions are needed to calculate gradients, e.g., sigmoid:

$$\psi_{\alpha}(v) = \frac{1}{1 + \exp(-\alpha v)}$$



# Backpropagation

◆ Derivative of a sigmoid function ( $\alpha = 1$ )


$$\nabla_v \psi(v) = \frac{e^{-v}}{(1 + e^{-v})^2} = \psi(v)(1 - \psi(v))$$

- Note about the small scale of the gradient
- Gradient vanishing issue

◆ Many other activation functions examined


# Gradient computation at output layer

- ◆ Output neurons are separate:



# Gradient computation at output layer

◆ Signal flow:



$$R_j = \frac{1}{2} e_j^2$$


$$R = \frac{1}{2} \sum_j e_j^2$$

$$\begin{aligned} \nabla_{w_{ji}} R &= \frac{\partial R_j}{\partial e_j} \frac{\partial e_j}{\partial d_j} \frac{\partial d_j}{\partial v_j} \frac{\partial v_j}{\partial w_{ji}} \\ &= e_j \cdot (-1) \cdot \psi'(v_j) \cdot f_i(\mathbf{x}) \\ &= -e_j \psi'(v_j) f_i(\mathbf{x}) \end{aligned}$$

Local gradient:  $\delta_j = -\frac{\partial R}{\partial v_j}$

# Gradient computation at hidden layer


- ◆ Output neurons are NOT separate:







# Gradient computation at hidden layer



$$v_i = (\mathbf{w}'_i)^\top \mathbf{g} \quad f_i = \psi(v_i) \quad v_j = \mathbf{w}_j^\top \mathbf{f} \quad d_j = \psi(v_j) \quad e_j = y_j - d_j$$

$$\begin{aligned}
 R_j &= \frac{1}{2} e_j^2 \\
 R &= \frac{1}{2} \sum_j e_j^2 \\
 \nabla_{w'_{ik}} R &= \sum_j \frac{\partial R_j}{\partial e_j} \frac{\partial e_j}{\partial d_j} \frac{\partial d_j}{\partial v_j} \frac{\partial v_j}{\partial f_i} \frac{\partial f_i}{\partial v_i} \frac{\partial v_i}{\partial w'_{ik}} \\
 &= - \sum_j e_j \psi'(v_j) w_{ji} \psi'(v_i) g_k(\mathbf{x}) \\
 &= - \sum_j \delta_j w_{ji} \psi'(v_i) g_k(\mathbf{x}) \quad \text{Local gradient: } \delta_i = - \frac{\partial R}{\partial v_i}
 \end{aligned}$$

# Back-propagation formula

◆ The update rule of **local gradients**:


□ for hidden neuron  $i$ :

$$\delta_i = \psi'(v_i) \sum_j \delta_j w_{ji}$$

↑

Only depends on the activation function at hidden neuron  $i$

◆ Flow of error signal:



# Back-propagation formula

◆ The update rule of weights:

□ Output neuron:

$$\Delta w_{ji} = \lambda \cdot \delta_j \cdot f_i(\mathbf{x})$$

□ Hidden neuron:

$$\Delta w'_{ik} = \lambda \cdot \delta_i \cdot g_k(\mathbf{x})$$

$$\begin{pmatrix} \text{Weight} \\ \text{correction} \\ \Delta w_{ji} \end{pmatrix} = \begin{pmatrix} \text{learning} \\ \text{rate} \\ \lambda \end{pmatrix} \cdot \begin{pmatrix} \text{local} \\ \text{gradient} \\ \delta_j \end{pmatrix} \cdot \begin{pmatrix} \text{input signal} \\ \text{of neuron } j \\ v_i \end{pmatrix}$$

# Two Passes of Computation

## ◆ Forward pass

- Weights fixed
- Start at the first hidden layer
- Compute the output of each neuron
- End at output layer

## ◆ Backward pass

- Start at the output layer
- Pass error signal backward through the network
- Compute local gradients

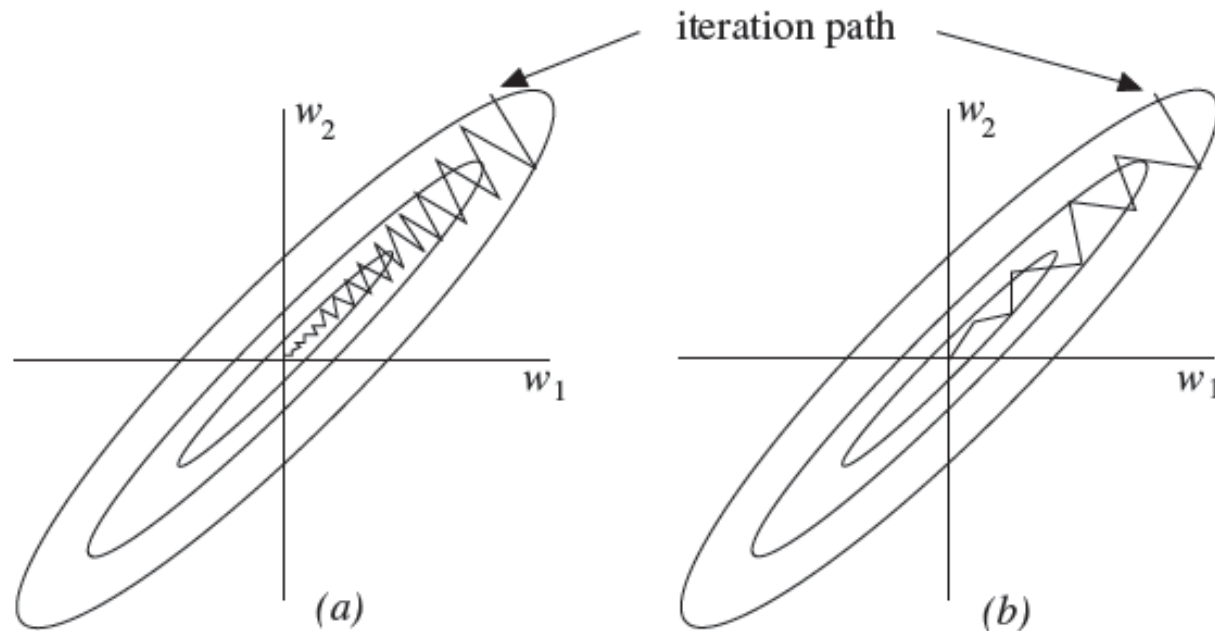
# Stopping Criterion

- ◆ No general rules
  
- ◆ Some reasonable heuristics:
  - The norm of gradient is small enough
  - The number of iterations is larger than a threshold
  - The training error is stable
  - ...

# Improve Backpropagation

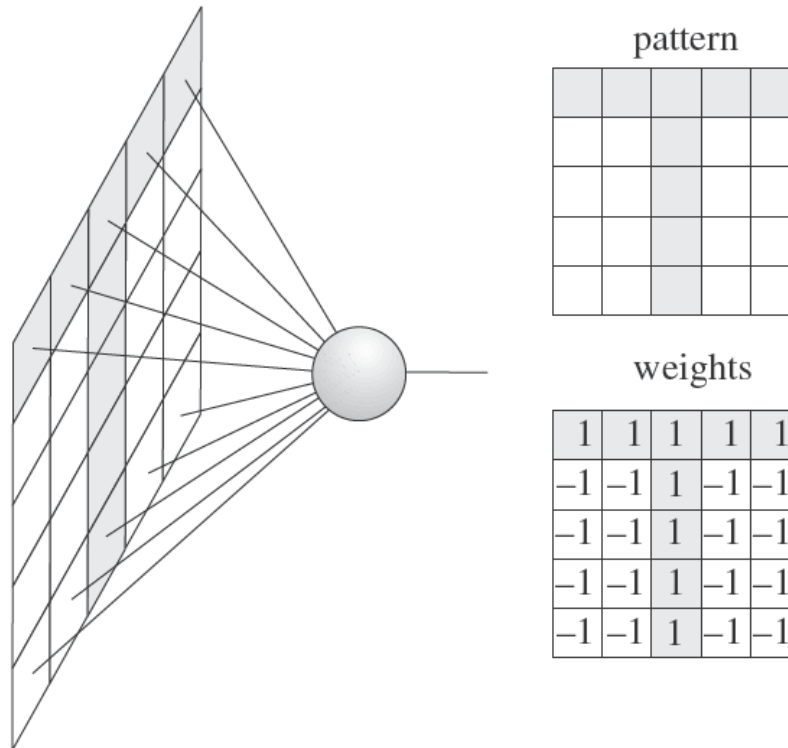
- ◆ Many methods exist to improve backpropagation
- ◆ E.g., backpropagation with momentum

$$\Delta w_{ij}^t = -\lambda \frac{\partial R}{\partial w_{ij}} + \alpha \Delta w_{ij}^{t-1}$$



# Neurons as Feature Extractor

- ◆ Compute the similarity of a pattern to the ideal pattern of a neuron
- ◆ Threshold is the minimal similarity required for a pattern
- ◆ Reversely, it visualizes the connections of a neuron



# Vanishing gradient problem

◆ The gradient can decrease exponentially during back-prop

◆ Solutions:

- Pre-training + fine tuning
- Rectifier neurons (sparse gradients)


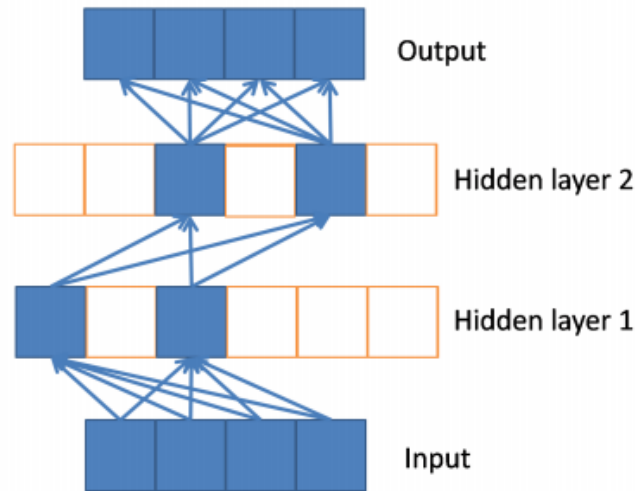
◆ Ref:

- Gradient flow in recurrent nets: the difficulty of learning long-term dependencies. Hochreiter, Bengio, & Frasconi, 2001



# Deep Rectifier Nets

- ◆ Sparse representations without gradient vanishing






- Non-linearity comes from the path selection
  - Only a subset of neurons are active for a given input
- Can be seen as a model with an exponential number of linear models that share weights

[Deep sparse rectifier neural networks. Glorot, Bordes, & Bengio, 2011]


# CNN

- ◆ Hubel and Wiesel's study on animal's visual cortex:
  - ❑ Cells that are sensitive to small sub-regions of the visual field, called a *receptive field*
  - ❑ Simple cells respond maximally to specific edge-like patterns within their receptive field. Complex cells have larger receptive fields and are locally invariant to the exact position of the pattern.




# Convolutional Neural Networks

- ◆ Sparse local connections (spatially contiguous receptive fields)



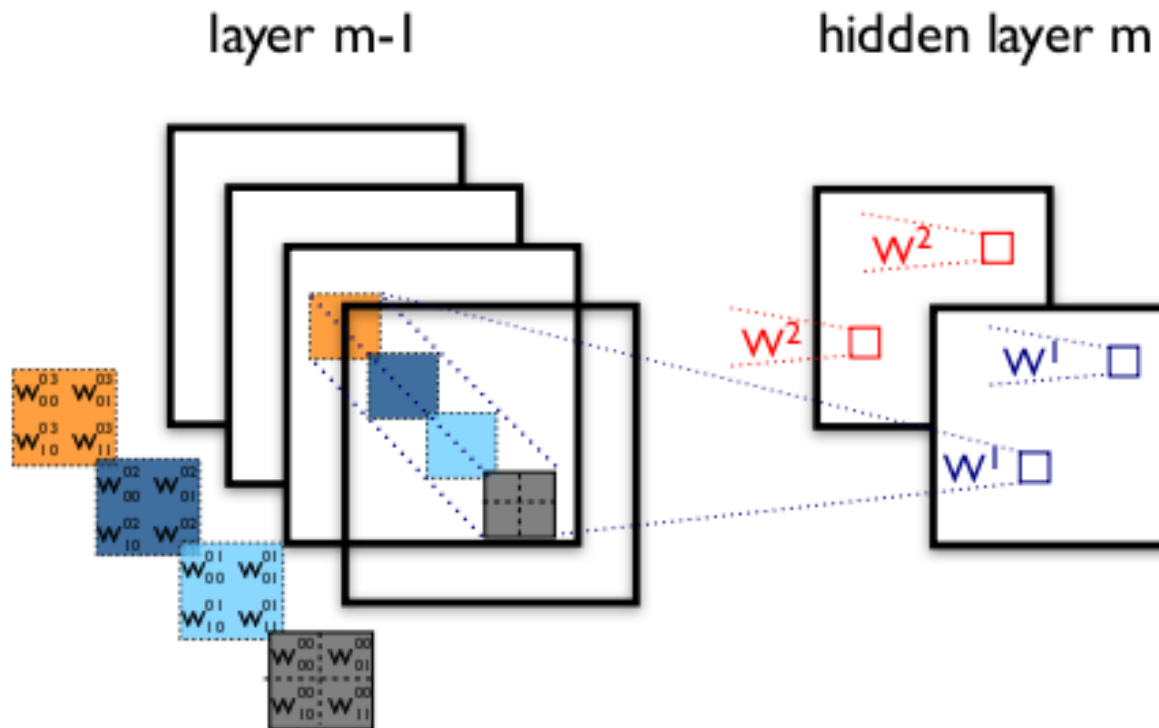
- ◆ Shared weights: each filter is replicated across the entire visual field, forming a feature map





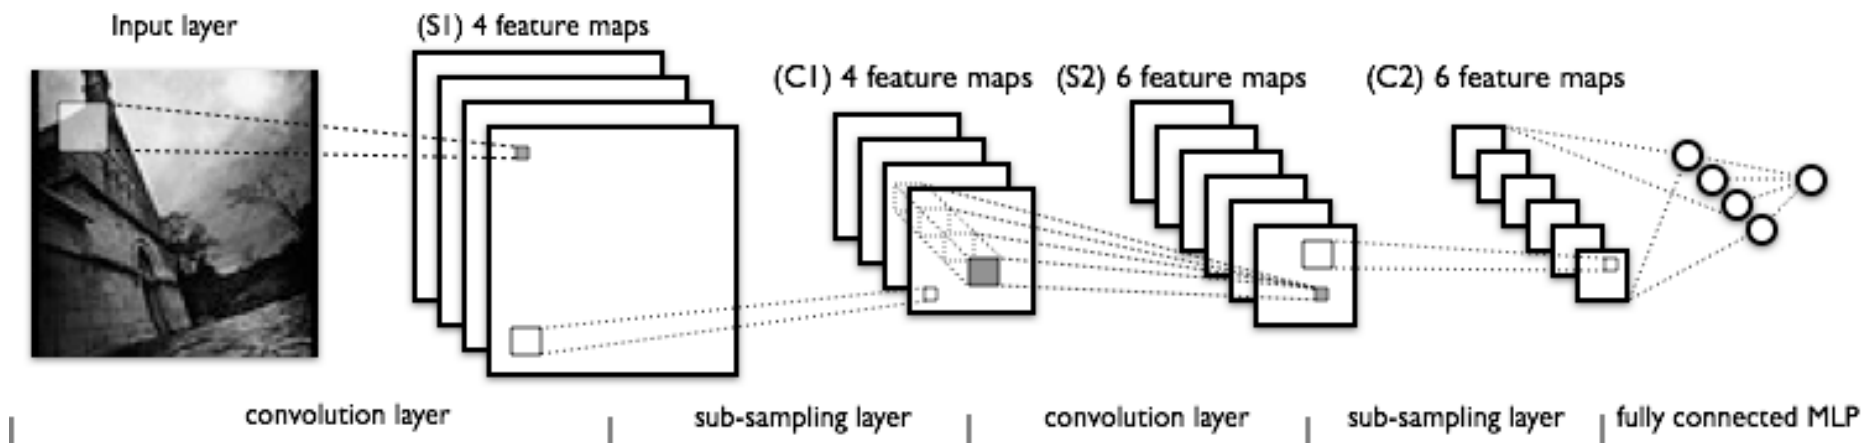
# CNN

- ◆ Each layer has multiple feature maps



# CNN

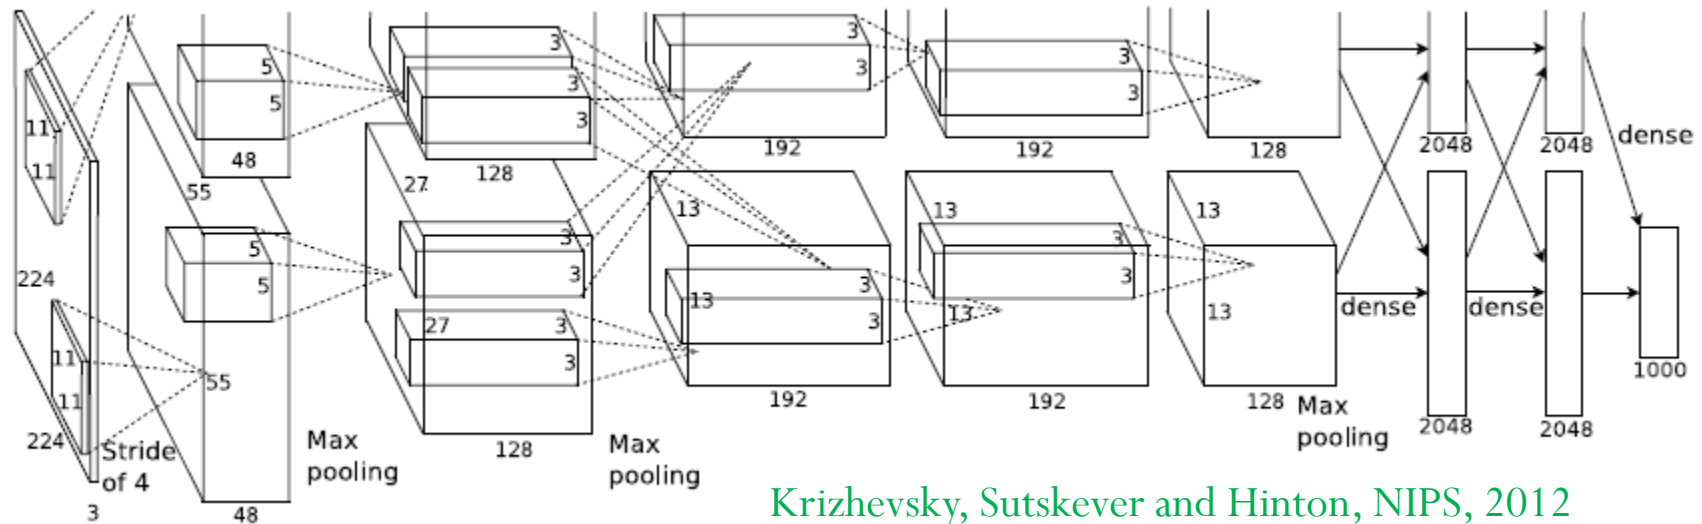
## ◆ The full model



## ◆ *Max-pooling*, a form of non-linear down-sampling.

- Max-pooling partitions the input image into a set of non-overlapping rectangles and, for each such sub-region, outputs the maximum value.

# Example: CNN for image classification



Krizhevsky, Sutskever and Hinton, NIPS, 2012


- ◆ Network dimension: 150,528(input)-253,440-186,624-64,896-64,896-43,264-4096-4096-1000(output)
  - In total: 60 million parameters
  - Task: classify 1.2 million high-resolution images in the ImageNet LSVRC-2010 contest into the 1000 different classes
  - Results: state-of-the-art accuracy on ImageNet

# Issues with CNN

- ◆ Computing the activations of a single convolutional filter is much more expensive than with traditional MLPs
- ◆ Many tuning parameters
  - # of filters:
    - Model complexity issue (overfitting vs underfitting)
  - Filter shape:
    - the right level of “granularity” in order to create abstractions at the proper scale, given a particular dataset
    - Usually 5x5 for MNIST at 1<sup>st</sup> layer
  - Max-pooling shape:
    - typical: 2x2; maybe 4x4 for large images

# Long Short-Term Memory

- ◆ A RNN architecture without gradient vanishing issue
- ◆ A RNN with LSTM blocks
  - Each block is a “smart” network, determining when to remember, when to continue to remember or forget, and when to output





## Discussions

# Challenges of DL

## ◆ Learning

- Backpropagation is slow and prone to gradient vanishing
- Issues with non-convex optimization in high-dimensions

## ◆ Overfitting


- Big models are lacking of statistical information to fit

## ◆ Interpretation

- Deep nets are often used as black-box tools for learning and inference

# Overfitting in DL

- ◆ Increasing research attention, e.g., dropout training (Hinton, 2012)




- ◆ More theoretical understanding and extensions
  - MCF (van der Maaten et al., 2013); Logistic-loss (Wager et al., 2013); Dropout SVM (Chen, Zhu et al., 2014)




# Some counter-intuitive properties

- ◆ Stability w.r.t small perturbations to inputs
  - Imperceptible non-random perturbation can arbitrarily change the prediction (**adversarial examples exist!**)



(a)

10x of  
differences



(b)

# Criticisms of DL

- ◆ Just a buzzword, or largely a rebranding of neural networks
- ◆ Lack of theory
  - gradient descent has been understood for a while
  - DL is often used as black-box
- ◆ DL is only part of the larger challenge of building intelligent machines, still lacking of:
  - causal relationships
  - logic inferences
  - integrating abstract knowledge



# Will DL make other ML methods obsolete?

**Quora** 2014/12/23


**Yes** (2 post, 113 upvotes)

- best predictive power when data sufficient
- DL is far from saturated
- Google et al invests on DL, it is the “richest” AI topic

**No** (10 posts, 284 upvotes)

- simpler algorithms are just fine in many cases
- methods with domain knowledge works better
- DL is feature learning, needs other methods to work
- DL is not that well developed, a lot of work to be done using more traditional methods
- No free lunch
- a lot like how ANN was viewed in the late 80s

Subjective Popularity



# What are people saying?

## ◆ Yann LeCun:

- “AI has gone from failure to failure, with bits of progress. This could be another leapfrog”

## ◆ Jitendra Malik:

- in the long term, deep learning may not win the day; ... “Over time people will decide what works best in different domains.”
- “Neural nets were always a delicate art to manage. There is some black magic involved”

## ◆ Andrew Ng:

- “Deep learning happens to have the property that if you feed it more data it gets better and better,”
- “Deep-learning algorithms aren't the only ones like that, but they're arguably the best — certainly the easiest. That's why it has huge promise for the future.”

# What are people saying?

## ◆ Oren Etzioni:

- “It’s like when we invented flight” (not using the brain for inspiration)

## ◆ Alternatives:

- Logic, knowledge base, grammars?
- Quantum AI/ML?



Thank You!