

CS838-1 Advanced NLP: Text Categorization with Naive Bayes Classifiers

Xiaojin Zhu

2007

Send comments to jerryzhu@cs.wisc.edu

Say you want to file your emails into two folders: “study”, “fun” (not that they should be different!). And you want to do it *automatically*. How?

1 Machine Learning Basics

- example, item, point, instance: each input object (e.g., a document, an image, a person).
- feature $x \in \mathcal{X}$: usually a fixed-dimensional numerical vector that characterizes the instance. For example, for a document x can be the word count vector. \mathcal{X} is the feature space.
- label $y \in \mathcal{Y}$: numerical encoding of output. It can be a discrete number (e.g. -1, 1 or 0, 1) for binary classification, $1, \dots, K$ for multiclass classification, a real number for regression.
- classifier: $f : \mathcal{X} \rightarrow \mathcal{Y}$ when \mathcal{Y} encodes discrete classes. This also implies that the particular class encoding is not important.
- training set: $\{(x_1, y_1), \dots, (x_n, y_n)\} \sim p(x, y)$. We assume the training examples are drawn i.i.d. from an unknown but fixed joint probability $p(x, y)$.
- training error (rate): $\frac{1}{n} \sum_{i=1}^n [f(x_i) \neq y_i]$.
- generalization error: $\mathbb{E}_p[f(x) \neq y]$. *Goal of machine learning*: Given training set, find f to minimize generalization error. This is difficult because we assume $p(x, y)$ is unknown.
- test set error: $\frac{1}{m} \sum_{i=n+1}^{n+m} [f(x_i) \neq y_i]$. The examples are again drawn i.i.d. from $p(x, y)$. Test set error is an estimate of generalization error.
- Overfitting: A classifier trained to minimize training error will often perform poorly on test set. This is known as overfitting. However, it is not

a good idea to *tune parameters to minimize test set error* either, because it is essentially overfitting the test set, and the true generalization error will be higher. Tuning on test set in most cases is regarded as *cheating* in machine learning.

- **Tuning set:** One may randomly split the training data into two parts: a smaller ‘training set’ (say 70%) and a ‘tuning set’ (say 30%). One trains a classifier from the training set, but tunes any parameters to minimize the *tuning set error* instead of training set error. The best model’s performance is then measure on the test set. This is a valid procedure because the test set is not used to select a classifier.
- **k -fold cross validation:** Sometimes one only has the training set $\{(x, y)_{1:n}\}$, but not a separate test set. One can simulate test set error as follows: Randomly split the training set into k equal *folds*. First, use folds $1, \dots, k-1$ to train a classifier, and treat fold k as the test set to compute error e_k . Second, use folds $1, \dots, k-2, k$ to train a different classifier, and treat fold $k-1$ as the test set to compute error e_{k-1} . Repeat the procedure for all k folds. Finally the *k -fold cross validation error* is the average of e_1, \dots, e_k . In order not to waste data, the final classifier is trained on the complete training set. When $k = n$, this is known as leave-one-out cross validation.

2 Naive Bayes Classifier

Let each document be represented by $x = (c_1, \dots, c_v)^\top$ the word count vector, otherwise known as *bag of word* representation. We assume within each class y , the probability of a document follows the multinomial distribution with parameter θ_y :

$$p(x|y) \propto \prod_{w=1}^v \theta_{yw}^{c_w}. \quad (1)$$

The log likelihood is

$$\log p(x|y) = x^\top \log \theta_y + \text{const}. \quad (2)$$

Note different classes have different θ_y ’s. Also note that the multinomial distribution assume *conditional independence* of feature dimensions $1, \dots, v$ given the class y . We know this is not true in reality, and more sophisticated models would assume otherwise. For this reason, such assumption on independence of features is known as the *naïve Bayes* assumption¹.

¹Whether you put two dots above i is a matter of personal taste.

Classification is done via Bayes rule:

$$y^* = \arg \max_y p(y|x) \quad (3)$$

$$= \arg \max_y \frac{p(x|y)p(y)}{p(x)} \quad (4)$$

$$= \arg \max_y p(x|y)p(y) \quad (5)$$

$$= \arg \max_y x^\top \log \theta_y + \log p(y), \quad (6)$$

where $p(y)$ is often estimated from the frequency of class y in training data. In this process we computed $p(y|x)^2$, and we assumed that the parameters θ_y are known for all classes. The process of computing the marginal distribution of unknown variable (y) given observed variables (x) is called *inference*.

Given a training set $\{(x_1, y_1), \dots, (x_n, y_n)\}$, *training* or *parameter learning* involves finding the best parameters $\Theta = \{\pi, \theta_1, \dots, \theta_C\}$. The model is $p(y = j) = \pi_j$, and $p(x|y = j) = \text{Mult}(x; \theta_j) \propto \prod_{w=1}^V \theta_{jw}^{x_{jw}}$. For simplicity we use the MLE here, but MAP is common too. We maximize the joint (log) likelihood of the training set:

$$\ell = \log p((x, y)_{1:n} | \Theta) ; \text{ hide } \Theta \text{ below} \quad (7)$$

$$= \log \prod_{i=1}^n p(x_i, y_i) \quad (8)$$

$$= \sum_{i=1}^n \log p(x_i, y_i) \quad (9)$$

$$= \sum_{i=1}^n \log p(y_i) + \log p(x_i | y_i). \quad (10)$$

We can formulate this as a constrained optimization problem,

$$\max_{\Theta} \quad \ell \quad (11)$$

$$\text{s.t.} \quad \sum_{j=1}^C \pi_j = 1, \quad C \text{ is the number of classes} \quad (12)$$

$$\sum_{w=1}^V \theta_{jw} = 1, \quad \forall j = 1 \dots C. \quad (13)$$

It is easy to solve it using Lagrange multipliers and arrive at

$$\pi_j = \frac{\sum_{i=1}^n [y_i = j]}{n} \quad (14)$$

$$\theta_{jw} = \frac{\sum_{i: y_i = j} x_{iw}}{\sum_{i: y_i = j} \sum_{u=1}^V x_{iu}}. \quad (15)$$

²We did not normalize it, but normalization could be done if desired.

2.1 Naive Bayes as a Generative Model

A generative model is a probabilistic model which describe the full generation process of the data, or the joint probability $p(x, y)$. Our Naive Bayes model consists of $p(y)$ and $p(x|y)$, which do just that: One can generate data (x, y) by first sample $y \sim p(y)$, and then sample word counts from the multinomial $p(x|y)$.

There is another family of models known as discriminative models, which do not model $p(x, y)$. Instead, they focuses on the conditional $p(y|x)$, or a similar but non-probabilistic quantity, which is directly related to classification. We will see our first discriminative model when we discuss logistic regression.

2.2 Naive Bayes as a Linear Classifier

Consider binary classification where $y = 0$ or 1 . Our classification rule with arg max can equivalently be expressed with log odds ratio

$$f(x) = \log \frac{p(y=1|x)}{p(y=0|x)} \quad (16)$$

$$= \log p(y=1|x) - \log p(y=0|x) \quad (17)$$

$$= (\log \theta_1 - \log \theta_0)^\top x + (\log p(y=1) - \log p(y=0)). \quad (18)$$

The decision rule is to classify x with $y = 1$ if $f(x) > 0$, and $y = 0$ otherwise. Note for given parameters, this is a *linear function* in x . That is to say, the Naive Bayes classifier induces a linear decision boundary in feature space \mathcal{X} . The boundary takes the form of a hyperplane.

2.3 Naive Bayes as a Special Case of Bayes Networks

A *Bayes Network* is a directed graph that represent a family of probability distributions. This is covered in detail in [cB] Chapter 8.1, 8.2. Outline:

- nodes: each node is a random variable. We have one y node, and v x_w nodes.
- directed edges: No directed cycles allowed, i.e. must be a DAG. For naive Bayes, from y to x_w .
- meaning: the joint probability on all nodes $s_{1:K}$ is factorized in a particularly form

$$p(s) = \prod_{i=1}^K p(s_i | \text{pa}(s_i)), \quad (19)$$

where $\text{pa}(s_i)$ are the parents of s_i . For naive Bayes, $p(x_{1:v}, y) = p(y) \prod_{i=1}^v p(x_i | y)$.

- observed nodes: nodes with known values, e.g. $x_{1:v}$. Shaded.
- plate: a lazy way to duplicate the node (and associated edges) multiple times. Our $x_{1:v}$ can be condensed into a plate.